stable-Lipchitz SCMs: D-separation and Cyclic Causality

David Reber
Columbia University
david.reber@columbia.edu

Abstract: Despite being the natural modeling choice for phenomena involving feedback, cyclic causality has seen limited use because many of the convienent guarentees of acylic SCMs fail to hold for general cyclic SCMs. One well-known hurdle to this modeling option is that d-separation fails to hold for general cyclic SCMs.
In this report, we present research about the validity of d-separation (aka. the dGMP: directed global Markov property) for cyclic SCMs. Specifically, we propose the class of stable-Lipschitz SCMs, which generalize acyclic SCMs and are contained by simple SCMs. Stable-Lipschitz SCMs behave like linear SCMs asyptotically, and indeed, are proven to satsify the dGMP (currently with some additional constraints). These results are verified numerically, and the validity of the backdoor adjustment criterion is proven as well. We further show that stable-Lipschitz SCMs are closed under interventions, and that the interventional distributions of simple SCMs satisfy the dGMP.
Lastly, we discuss future directions for research, including the Pearl Causal Hierarchy, the do-calculus, and most exciting, the possibility of causal identification with multiple equilibria (e.g. in game theory, economics).

1. Introduction

1.1. Motivation and Problem Formulation

To the extent that feedback is present in a system (i.e. game theory, economics, neuroscience, the physical sciences generally) a natural modeling choice would be to represent this feedback as a cycle in the structural causal model; after all, there is a reason dynamical models have been long dominant for explaining these phenomenon. To stubbornly try to fit acyclic models to known cyclic phenomena seems, at best, akin doing science with one arm tied behind one's back; at worst, it amounts to fitting epicycles.

However, there are strong reasons to prefer having an arm tied behind one's back: the introduction of cycles breaks many of the convienent guarantees of acyclic models [1]: the lack of a unique equilibrium means the potential response function is not well-defined; the observational, interventional, and counterfactual distributions may not exist, or if they do, they may not be unique; marginalizing over variables may not be possible, and even if it is, the causal semantics may not be preserved; d-separation may not hold (aka. the "directed global Markov property"); even the weaker variant of σ -separation may not hold (the "general directed global Markov property").

Consequently, the 'best of both worlds' would be to find some space of SCMs which is at least somewhat more general than acyclic SCMs, while preseving their most helpful properties. [1] presents *simple SCMs* as a such a class; however, while σ -separation holds for this class, d-separation in general does not. Conseqently, I propose a new class of SCMs, *stable-Lipschitz SCMs*, which satisfies the directed global Markov property. This class is properly contained within simple SCMs, and as an added bonus, provides a simple way to verity that an SCM is indeed simple. See Figure 1 for a visualization of this set inclusion.

Formally, we seek to answer the following problem:

Problem 1 (stable-Lipschitz SCMs and d-separation: Observational)**.** *Prove whether the observational distribution of stable-Lipschitz SCMs satisfy the directed global Markov property: that is, whether every conditional independence read-off by d-separation in the causal graph G_M holds in the observational distribution $P_M(V)$.*
Inputs/Outputs: Expressed by the green implication in Figure 2.

After showing that Problem 1 resolves positively, we then address the following problem as well:

Problem 2 (stable-Lipschitz SCMs and d-separation: Interventional)**.** *Prove whether the interventional distribution of stable-Lipschitz SCMs satisfy the directed global Markov property: that is, whether every conditional independence read-off by d-separation in $G_{\overline{\mathbf{X}}}$ holds in the interventional distribution $P_{M_{do(\mathbf{x})}}(\mathbf{V})$.*

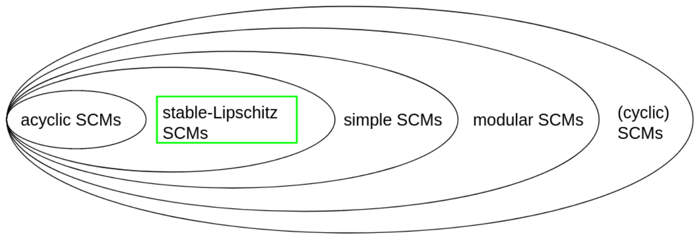

Figure 1: Proven set-inclusion relation of the class of stable-Lipschitz SCMs relative to other classes of cyclic SCMs. Green is my contribution, the rest is from [1]

$$\begin{tikzcd} & P_M(V) \arrow[r] & (X \perp\!\!\!\perp Y|Z)_{P_M} \\ M \arrow[ur] \arrow[dr] & & \\ & G_M \arrow[r] & (X \perp\!\!\!\perp Y|Z)_{G_M} \arrow[uu, Rightarrow] \end{tikzcd}$$

Figure 2: When M is solvable, the green implication is the directed global Markov property (validity of d-separation).

1.2. Literature Review and Framework Selection

The attractiveness of incorperating feedback into causality has already attracted significant intellectual effort, although a clear consensous has not yet been reached regarding what mathematical formulation to use. Perhaps the most intuitive approach to those familiar with DAG-based causality, is to unravel the feedback by duplicating the endogenous variables and indexing w.r.t time. However, as [10] points out, this model breaks down if the rate at which data is sampled is significantly slower than the feedback frequency.

Settable systems are a well-known framework for cyclic causality [14] which I considered using, as it has proven insightful for game theory [15]. The primary strength of settable systems lies in its convenient treatment of multiple equilibria. However, since stable-Lipschitz SCMs do not have multiple equilibria, this is more functionality than is needed right now. Settable systems also rely on fairly low-level representations of the optimization pressure on the equilibria (for example, to select one equilibria from many) [8], and the addition of additional variables for distinguising between equilibria is more of a departure from [11] than is currently needed for this research.

Meanwhile, the modeling framework for *cyclic SCMs* of [1] is, at least in spirit, very similar to Pearl's framework for acyclic SCMs [11]. Cycles are incorporated simply by allowing the structural function F of the SCM $M = (U,V,F,P(U))$ to be cyclic, that is, dependent on any variables in $U \cup V$. This framework is currently developed for atomic interventions (which the authors rename 'perfect interventions'). The authors take a more abstract approach than either of the two alternatives discussed above, for better and for worse. Consequently, this framework makes it straightforward to show the inheritance of properties between various classes of SCMs.

However, there are several aspects of this framework which contrast with Pearl [11]. The authors emphasize almost-everywhere equivalence, with respect to the observational, interventional, and counterfactual distributions, or structural functions themselves. This emphasis implicitly impacts their definitions. For example, and notably, the definition of a variable's 'parents' excludes any trivial dependencies. While this definition is semantically useful (we only say there's a parent if there's a dependence that really can't be ignored) it can be slightly confusing: for example, consider the two-node linear system of $x_1 = f_1(x_1,x_2) = 0.5x_1 + 0.1x_2$ and $x_2 = f_2(x_1,x_2) = 0.1x_1 + 0.5x_2$. According to the definitions of [1], x_1 and x_2 are both parents of each other, yet no self-cycles exist, because there exists an SCM which is asymptotically equivalent without self-cycles. (This property of linear systems is discussed further in [10]).

Despite these cumbersome definitions, I have opted to follow this framework for now, because 1. the inheritance of properties makes proofs more clear, 2. it seems overall the most similar to Pearl's framework in [11], and 3. it has also been shown applicable to modeling game theory [8].

Lastly, I would like to highlight the primary contribution by [1] of *simple SCMs*, as this class serves both as inspiration for and contrast to the class of *stable-Lipschitz SCMs*. The authors demonstrate that simple SCMs preserve many of the properties of acyclic SCMs (although notably, not d-separation). Simple SCMs are defined as SCMs which are uniquely solvable with respect to every subset of variables. Consequently, the observational, interventional, and counterfactual distributions exist and are unique. Furthermore, simple SCMs satisfy the general directed global Markov property; are closed under marginalization; and the causal semantics of the causal

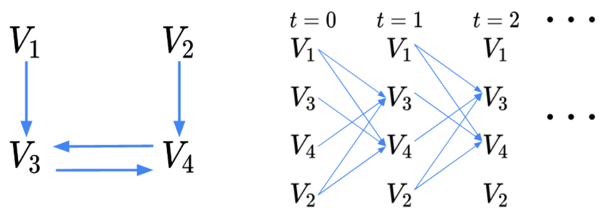

Figure 3: Left: The cyclic graph corresponding to the structural equations of Examples 2 and 3. Right: Rolling out this graph to explicitly represent the transient dynamics.

SCM	d-separation valid?
Acylic	Yes
Discrete	Yes
Linear	Yes
Nonlinear, Continuous	Not in general (non-Lipschitz counterexample)
Lipschitz, Continuous	????

Figure 4: Classes of SCMs known to satisfy the directed global Markov property, as summarized by [1]. Here *discrete* and *continuous* refer to the domains of the exogenous and endogenous variables.

diagram's subgraph match the structure of the submodel of the SCM.

1.3. Previous Results regarding D-separation

When d-separation is not valid (the directed global Markov property does not hold), it may be that a weaker condition is satisfied. σ-separation is an extension of d-separation, which works by applying d-separation to the acyclification of the original graph. σ-separation implies d-separation; in other words, the general directed global Markov property is weaker than the directed global Markov property. The relationship between these two Markov properties can be seen in Figure 5.

Acyclic SCMs, discrete SCMs (with ancestrally unique solvability), and linear SCMs (with nontrivial dependencies and positive measure) are known to satisfy the directed global Markov property [7]. On the other hand, simple SCMs are known to satisfy the general directed global Markov property [1]. These results are summarized in Figure 4.

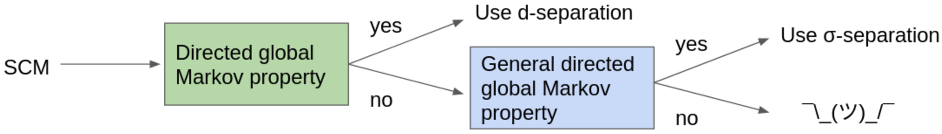

Figure 5: Relationship between d-separation and σ-separation.

1.4. Summary of Main Results

The primary contributions of this paper are as follows, with the shorthand of dGMP := Directed Global Markov Property (a.k.a. "Validity of d-separation").

- Theorems
  - acyclic SCMs $\subset$ stable-Lipschitz SCMs $\subset$ simple SCMs
  - stable-Lipschitz SCMs are closed under interventions
  - **Obs. D-separation**: The observational distribution of stable-Lipschitz SCMs satisfies dGMP
  - Validity of Backdoor Criterion on stable-Lipschitz SCMs
  - **Int. D-separation**: The interventional distribution of stable-Lipschitz SCMs satisfies dGMP
  - stable-Lipschitz SCMs are closed under twin-operation
- Numerics
  - Obs. D-separation (above)
  - The observational distribution of (non-stable) Lipschitz SCMs also satisfy dGMP

1.5. Paper Organization

This paper is organized as follows. First, Section 2 discusses the intution of the relationship between dynamics and potential response functions, by exploring three examples of increasing complexity. Next in Section 3, we examine a case study for how d-separation can fail, from several perspectives, and how it can be modified to be stable-Lipschitz. That this example satisfies d-separation sets up the main theorems of Section 4. Numerical investigations of my conjectures are presented in Section 5. Section 6 discusses the next steps of the research.

2. Dynamics and Potential Responses

Let's examine how the dynamics and potential response functions change as we increase the complexity of the SCM, while introducing some important definitions along the way:

- acyclic, linear SCMs (Example 1)
- cyclic, linear SCMs (Example 2)
- cyclic, nonlinear SCMs (Example 3)

2.1. Acyclic, Linear

Acyclic SCMs have the property that once the values of early nodes are fixed, later nodes in the topological ordering are uniquely determined. For causal models with feedback to have this same behavior, there must exist unique solutions to subsets of the structural equations.

Definition 1 (Unique Solvability)**.** *Let $M = \langle V, U, F, P(U) \rangle$ be an SCM and $\mathbf{Z} \subseteq \mathbf{V}$ a subset of endogenous variables. We say that M is uniquely solvable w.r.t. $\mathbf{Z}$ if for almost every $\mathbf{u} \in dom(\mathbf{U})$ and $\mathbf{v}_{\setminus \mathbf{Z}} \in dom(\mathbf{V}_{\setminus \mathbf{Z}})$ the equations*

$$\mathbf{V_Z} = F_{\mathbf{Z}}(\mathbf{V}, \mathbf{u})$$

have a unique solution. If $\mathbf{Z} = \mathbf{V}$, we say that M is uniquely solvable.

([1] proves the equivalence of this condition with that of 'mapping fixed inputs to unique outputs', which will be made clear with Example 1).

Throughout this paper we only consider uniquely solvable SCMs, so we may use the following defintion of *potential response*.

Definition 2 (Potential Response)**.** *Let $M = \langle V, U, F, P(U) \rangle$ be a uniquely solvable SCM. The potential response function is the mapping $\overline{\mathbf{V}}(\mathbf{U}) : dom(\mathbf{U}) \to dom(\mathbf{V})$ which associates each $\mathbf{u} \in dom(\mathbf{U})$ with the unique solution $\mathbf{v}^*$ of $\mathbf{V} = F(\mathbf{V}, \mathbf{u})$.*

The following example solidifies these concepts.

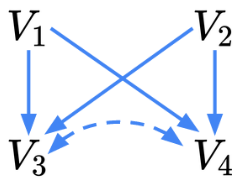

Figure 6: The causal graph of Example 1.

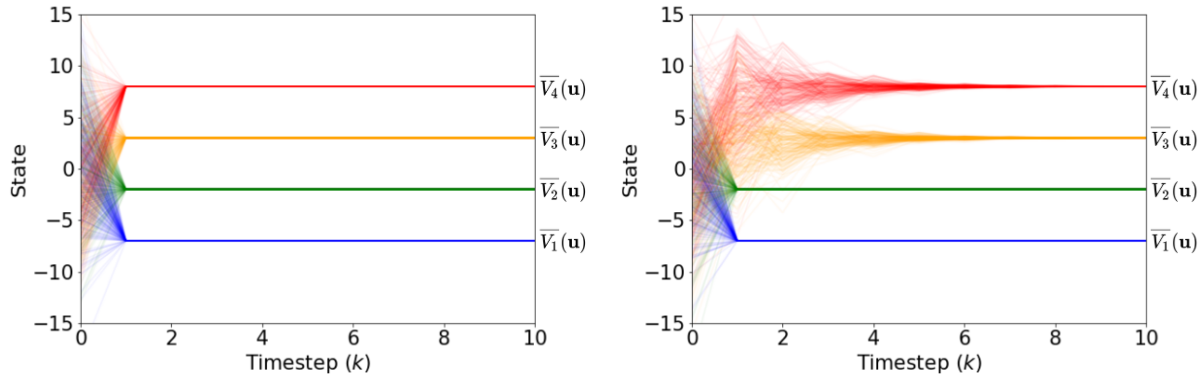

Figure 7: The dynamics of Example 1 (left) and Example 2 (right), with fixed $\mathbf{u} = (-7, -2, 2.5, 7.5)$, and initialization $\mathbf{v}_0 \sim \mathcal{N}(0, 5)$.

Example 1 (Acyclic, Linear SCM)**.** *Consider the following acyclic, linear SCM*

$$M_1 = \begin{cases} \mathbf{u}, \mathbf{v} \in \mathbb{R}^4, \quad U_i \sim \mathcal{N}(0,1) \quad iid. \\ f_1 : V_1 \leftarrow U_1 \\ f_2 : V_2 \leftarrow U_2 \\ f_3 : V_3 \leftarrow \frac{2}{3}V_1 + \frac{1}{3}V_2 + \frac{4}{3}U_3 + \frac{2}{3}U_4 \\ f_4 : V_4 \leftarrow \frac{1}{3}V_1 + \frac{2}{3}V_2 + \frac{2}{3}U_3 + \frac{4}{3}U_4 \end{cases}$$

where $\mathcal{N}(0,1)$ is Gaussian. The causal graph of M_1 is shown in Figure 6.

To evaluate the potential response $\overline{\mathbf{V}}(\mathbf{U})$ of M_1, we may sample $\mathbf{u} \in \mathbf{U}$ and consider the limit

$$\overline{\mathbf{V}}(\mathbf{U}) = \lim_{k \to \infty} F^{(k)}_{\mathbf{U}=\mathbf{u}}(\mathbf{V}_0)$$

In this example, the limit exists and is unique because M_1 is acyclic. In fact, the dynamics converge within a single timestep, as shown (for the particular $\mathbf{u} = (-7, -2, 2.5, 7.5)$) in Figure 7 (left).

In this case we can evaluate the potential response analytically for any $\mathbf{u} \in \mathbf{U}$ as

$$\begin{aligned} \overline{V_1}(\mathbf{U}) &= U_1 \\ \overline{V_2}(\mathbf{U}) &= U_2 \\ \overline{V_3}(\mathbf{U}) &= \frac{2}{3}U_1 + \frac{1}{3}U_2 + \frac{4}{3}U_3 + \frac{2}{3}U_4 \\ \overline{V_4}(\mathbf{U}) &= \frac{1}{3}U_1 + \frac{2}{3}U_2 + \frac{2}{3}U_3 + \frac{4}{3}U_4 \end{aligned}$$

Example 1 highlights why unique solvability is such a desirable property of cyclic SCMs: the choice of the initialization $\mathbf{v}_0$ is 'washed-out'. In general, cyclic SCMs may not be solvable at all, or the equilibria may not be unique; furthermore, these properties are not generally invariant under interventions or marginalizations [1].

2.2. *Cyclic, Linear*

Now, let us consider a cyclic, linear SCM, which is otherwise as close to Example 1 as possible.

Example 2 (Cyclic, Linear SCM)**.** *Consider*

$$M_2 = \begin{cases} \mathbf{u}, \mathbf{v} \in \mathbb{R}^4, \quad U_i \sim \mathcal{N}(0,1) \quad iid. \\ f_1 : V_1 \leftarrow U_1 \\ f_2 : V_2 \leftarrow U_2 \\ f_3 : V_3 \leftarrow \frac{1}{2}V_1 + \frac{1}{2}V_4 + U_3 \\ f_4 : V_4 \leftarrow \frac{1}{2}V_2 + \frac{1}{2}V_3 + U_4 \end{cases}$$

The causal graph of M_2 is shown in Figure 3 (left). Note that there are now two edges between V_3 and V_4, as each depends on the other.

As Figure 7 (right) shows, the feedback of the structural equations delays convergence of the endogenous variables to the fixed point. Note however, that nodes V_1 and V_2 precede the strongly connected component $\{V_3, V_4\}$ in the topological ordering, and converge immediately (as if they were acyclic).

By construction, M_1 and M_2 have the same potential response $\overline{\mathbf{V}}(\mathbf{U})$:

$$\begin{aligned} \overline{V_1}(\mathbf{U}) &= U_1 \\ \overline{V_2}(\mathbf{U}) &= U_2 \\ \overline{V_3}(\mathbf{U}) &= \frac{2}{3}U_1 + \frac{1}{3}U_2 + \frac{4}{3}U_3 + \frac{2}{3}U_4 \\ \overline{V_4}(\mathbf{U}) &= \frac{1}{3}U_1 + \frac{2}{3}U_2 + \frac{2}{3}U_3 + \frac{4}{3}U_4 \end{aligned}$$

Thus, M_1 and M_2 are observationally equivalent; that is, they induce the same observational distribution, even though one has feedback and the other does not. Note that in this case, modeling the feedback in M_2 produces a simpler structural form than M_1.

In general, if a cyclic SCM is linear, so is its potential response (if it exists). To see this, note that for linear structural equations $F(\mathbf{V}) = B\mathbf{V} + \mathbf{U}$ we can derive the general form of the potential response as $\overline{\mathbf{V}} = (I - B)^{-1}\mathbf{U}$ whenever $\det(I - B) \neq 0$.

For linear (cyclic) SCMs, the Lipschitz bound on the dynamics is particularly easy to evaluate, as it amounts to taking the entry-wise absolute-value of the transition matrix: $|A|$. More generally, we use the following defintion adopted from [12]:

Definition 3 (Lipschitz Matrix)**.** *Let $M = \langle V, U, F, P(U) \rangle$ be an SCM, with $dom(\mathbf{U}) = \mathbb{R}^m$, $dom(\mathbf{V}) = \mathbb{R}^n$, and $F : dom(\mathbf{U}) \times dom(\mathbf{V}) \to dom(\mathbf{V})$ differentiable and Lipschitz.*

Let $\mathbf{Z} \subseteq \mathbf{V}$ be a subset of endogenous variables. For each pair of vertex indices $i, j \in \mathbf{Z}$, define

$$a_{ij} = \sup_{\mathbf{u},\mathbf{v}} \left| \frac{\partial f_i}{\partial v_j}(\mathbf{u}, \mathbf{v}) \right|$$

We call the matrix $A_{\mathbf{Z}} = [a_{ij}]_{i,j \in \mathbf{Z}}$ the Lipschitz matrix of $F_{\mathbf{Z}}$.

When $\mathbf{Z} = \mathbf{V}$, we simply call $A = [a_{ij}]$ the Lipschitz matrix of F.

Consider again Example 2. The Lipschitz matrix of M_2 is

$$A = \begin{bmatrix} 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 \\ \frac{1}{2} & 0 & 0 & \frac{1}{2} \\ 0 & \frac{1}{2} & \frac{1}{2} & 0 \end{bmatrix}$$

Note that A is zero where there is no functional dependence, and otherwise represents the Lipschitz constant along the direction of each partial derivative.

The class of SCMs which are the focus of this paper are those whose strongly connected components have stable Lipschitz matricies (in order to ensure the unique solvability of the SCM).

Definition 4 (stable-Lipschitz SCM)**.** *Let $M = \langle V, U, F, P(U) \rangle$ be an SCM with causal diagram G. We say M is stable-Lipschitz if, for every strongly connected component $\mathbf{Z}$ of G, the following conditions hold:*

- *$F_{\mathbf{Z}}$ is differentiable and Lipschitz.*
- $\rho(A_{\mathbf{Z}}) < 1$.

where $\rho(A)$ is the spectral radius of A.

Again returning to Example 2, since the only strongly connected component in the causal graph (Figure 3) is $\{V_3, V_4\}$, we construct the Lipschitz matrix

$$A_{2,3} = \begin{bmatrix} \frac{1}{2} & \frac{1}{2} \\ \frac{1}{2} & \frac{1}{2} \end{bmatrix}$$

and check if the spectral radius of $A_{2,3}$ is less than 1. In this case, $\rho(A_{2,3}) = 1$ so M_2 is **not** *stable-Lipschitz*, even though it is *stable* (always converges to the unique potential response found in Example 2) and *Lipschitz*.

Remark 1 (Limitations of stable-Lipschitz)**.** *As Example 2 highlights, the definition of simple-Lipschitz is quite restrictive (as is highlighted later by the numerics). However, this class has convienent theoretical properties which I expect to be valuable for proving Conjecture 2.*

2.3. *Cyclic, Nonlinear*

Now we consider what the potential response function looks like, when the cyclic SCM in question is no longer linear (or even Lipschitz-continuous).

Example 3 (Cyclic, Nonlinear SCM [13])**.** *Consider*

$$M_3 = \begin{cases} \mathbf{u}, \mathbf{v} \in \mathbb{R}^4, \quad U_i \sim \mathcal{N}(0,1) \quad iid. \\ f_1 : V_1 \leftarrow U_1 \\ f_2 : V_2 \leftarrow U_2 \\ f_3 : V_3 \leftarrow V_1 V_4 + U_3 \\ f_4 : V_4 \leftarrow V_2 V_3 + U_4 \end{cases}$$

M_3 respects the same causal graph as M_2, shown in Figure 3 (left).

The potential response $\overline{\mathbf{V}}(\mathbf{U})$ of M_3 can be solved for analytically as

$$\begin{aligned} \overline{V_1}(\mathbf{U}) &= U_1 \\ \overline{V_2}(\mathbf{U}) &= U_2 \\ \overline{V_3}(\mathbf{U}) &= \frac{U_1 U_4 + U_3}{1 - U_1 U_2} \\ \overline{V_4}(\mathbf{U}) &= \frac{U_1 U_4 + U_3}{1 - U_1 U_2} U_2 + U_4 \end{aligned}$$

In particular, note the singularity when $U_1 U_2 = 1$. If V_3 and V_4 are observed to be large, then this means the exogenous variables are located near the singularity. This in turn implies that V_1 and V_2 are inversely proportional, so if V_1 is large, then V_2 is small. Consequently $(V_1 \not\perp\!\!\!\perp V_2 \mid V_3, V_4)_{P_{M_3}}$.

The Lipschitz matrix $A_{2,3}$ of the strongly connected component cannot be defined over this domain, since

$$a_{31} = \sup_{\mathbf{v} \in \mathbf{V}} |V_4| \quad a_{34} = \sup_{\mathbf{v} \in \mathbf{V}} |V_1| \quad a_{42} = \sup_{\mathbf{v} \in \mathbf{V}} |V_3| \quad a_{43} = \sup_{\mathbf{v} \in \mathbf{V}} |V_2|$$

are all unbounded over $\mathbf{u} \in \mathbb{R}^4$.

Thus, nonlinear cyclic SCMs can have singularities in their potential response functions, unlike linear SCMs. This consequently causes the Lipschitz matrix to be unbounded (and consequently not well-defined).

Remark 2. *As discussed in [13] (and further in [1]), Example 3 is a counterexample to the claim that d-separation works in general for nonlinear, continous domained cyclic SCMs. One way to see this (as pointed out by [13]) is to solve for the induced observational distribution directly:*

$$P(\mathbf{V}) = (\frac{1}{4\pi^2}) e^{-\frac{v_1^2}{2}} e^{-\frac{v_2^2}{2}} e^{-\frac{(v_3 - v_1 v_4)^2}{2}} e^{-\frac{(v_4 - v_2 v_3)^2}{2}} |(1 - v_1 v_2)^{-1}|$$

$P(\mathbf{V})$ cannot be factored completely into terms which not jointly contain v_1 and v_2

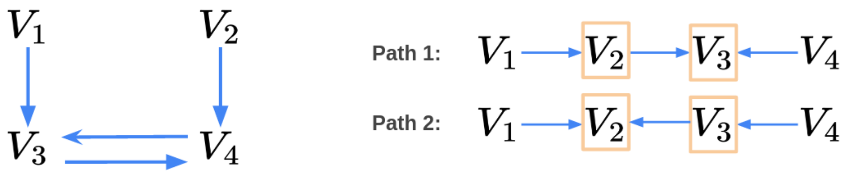

Figure 8: The paths which need to be considered to check $(V_1 \perp\!\!\!\perp V_2 \mid V_3, V_4)_G$.

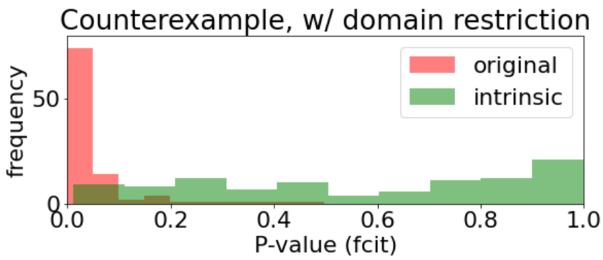

Figure 9: Numerical conditional tests of Examples 3 and 4. Here, conditional independence is tested while restricting the exogenous variables to be drawn from a truncated multivariate normal distribution (such that the SCM as a whole is intrinsically stable). The conditional independence identified by d-separation is now present in the observational distribution.

3. Cyclic D-separation

Surprisingly, even though we have $(V_1 \not\perp\!\!\!\perp V_2 \mid V_3, V_4)_{P_{M_3}}$ in the observational distribution of Example 3, we have $(V_1 \perp\!\!\!\perp V_2 \mid V_3, V_4)_G$ according to d-separation. To evaluate the statement, we consider both paths from V_1 to V_2 as shown in Figure 8. Since each of these paths is blocked by V_3, V_4, we have that $(V_1 \perp\!\!\!\perp V_2 \mid V_3, V_4)_G$. Thus, d-separation fails to be valid in this example.

Even though $(V_1 \perp\!\!\!\perp V_2 \mid V_3, V_4)_G$, $(V_1 \not\perp\!\!\!\perp V_2 \mid V_3, V_4)_{P_{M_3}}$ in Example 3, it seems that perhaps this failure of d-separation is due to the singluarity of the potential response function. In particular, note that the structural functions of Example 3 are locally-linear. Perhaps if we restrict the domain of M_3 to where it behaves like a linear SCM, it will still preserve d-separation?

Example 4 (Localization of an SCM)**.** *Consider*

$$M_4 = \begin{cases} \mathbf{u}, \mathbf{v} \in [-0.5, 0.5], \quad U_i \sim \mathcal{N}(0,1) \cap [-0.5, 0.5] \quad iid. \\ f_1 : V_1 \leftarrow U_1 \\ f_2 : V_2 \leftarrow U_2 \\ f_3 : V_3 \leftarrow V_1 V_4 + U_3 \\ f_4 : V_4 \leftarrow V_2 V_3 + U_4 \end{cases}$$

the only difference from M_3 being that we restrict the domain of the exogenous variables (and hence, the endogenous variables) to be closer to 0.

Since the structural functions did not change, the potential response function is the same. However, the Lipschitz matrix $A_{2,3}$ is now well defined

$$a_{31} = \sup_{\mathbf{v} \in \mathbf{V}} |V_4| = 0.5, \quad a_{34} = \sup_{\mathbf{v} \in \mathbf{V}} |V_1| = 0.5, \quad a_{42} = \sup_{\mathbf{v} \in \mathbf{V}} |V_3| = 0.5, \quad a_{43} = \sup_{\mathbf{v} \in \mathbf{V}} |V_2| = 0.5$$

and satifies $\rho(A_{2,3}) < 1$. Hence, M_4 is stable-Lipschitz.

Indeed, by bounding the domain of M_4 away from the sigularity, then d-separation works, as Figure 9 demonstrates numerically, as we would expect if Problem 1 resolves in the affirmative. In order to test the conditional independence numerically over a continuous domain, the python package `fcit` [6] was used. The resulting p-values are distributed roughly uniformly on $[0, 1]$, consistent with an `fcit` result of conditional independence.

4. Main Results

Note that in Definition 4 no constraints are placed on components of F which *are not* part of strongly connected components of G. This immediately implies the following reuslt:

Theorem 1 (acylic $\subset$ stable-Lipschitz)**.** *Let M be an acyclic SCM. Then M is stable-Lipschitz.*

Convienently, the restriction $\rho(A) < 1$ in Definition 4 is sufficient to ensure that stable-Lipschitz SCMs are uniquely solvable.

Lemma 2 (Unique Solvability)**.** *Let M be a stable-Lipschitz SCM. Then M is uniquely solvable.*

Simple SCMs are ones for which every submodel of the SCM is uniquely solvable. Recall from Definition 2 that this is necessary for the existence of a unique potential response function:

Definition 5 (Simple SCM [1])**.** *Let $M = \langle V, U, F, P(U) \rangle$ be an SCM. We call M simple if M is uniquely solvable w.r.t. every subset $\mathbf{Z} \subseteq \mathbf{V}$.*

Indeed, in [1] simple SCMs were proven to generalize a number of valuable properties about acyclic SCMs:

- the potential response function is always well-defined
- the observational, interventional, and counterfactual distributions always exist and are unique
- marginalizing over variables is always possible, and the causal semantics are always preserved
- σ -separation always holds (the "general directed global Markov property").

Stable-Lipschitz SCMs inherit all of these properties as well, because they are in fact contained within the space of simple SCMs:

Theorem 3 (stable-Lipschitz $\subset$ simple)**.** *Let M be a stable-Lipschitz SCM. Then M is simple.*

Together, these two inclusions give us Figure 1.
In addition, stable-SCMs are closed under interventions:

Theorem 4 (Closed under Interventions)**.** *Let M be a stable-Lipshitz SCM, $\mathbf{X} \subseteq \mathbf{V}$, and $\mathbf{x} \in dom(\mathbf{X})$. Then $M_{do(\mathbf{X}=\mathbf{x})}$ is stable-Lipschitz.*

This fact is used for demonstrating that d-separation holds both for the observational, and the interventional distributions.

Theorem 5 (Observational dGMP)**.** *Let M be stable Lipschitz with 1. structural equations of the form $F(\mathbf{V}, \mathbf{U}) = H(\mathbf{V}) + \mathbf{U}$ (additive noise), 2. each $U_i \cap U_j = \emptyset$ for $i \neq j$ (independent noise), and 3. $P_M(\mathbf{V})$ has density according the the Legesgue measure on $\mathbb{R}^{|\mathbf{V}|}$ (positivity).*
Then M satisfies the directed global Markov property.

I am confident that conditions 1 and 2 in the hypothesis can be substantially weakened with further research.
This is a very new result, so while I believe the proof to be accurate and comprehensive, I'm still vetting it for errors: I'd place 5:1 odds against finding an irrecoverable error in the proof.

Corollary 1 (Adjustment Formula)**.** *Let M be as in Theorem 5 and $Q = P(y|do(x))$ a causal query. If the BDC is satisfied, then Q can be found via backdoor adjustment.*

One of the motivations for weakening the condition of independent noise in Theorem 5 is to be able to similarly prove that the front-door criteria is also valid.

Theorem 6 (Interventional dGMP)**.** *Let M be as in Theorem 5. For any $\mathbf{X} \subseteq \mathbf{Z}$, $M_{do(\mathbf{X}=\mathbf{x})}$ satisfies the directed global Markov property.*

Because of their relation to *intrinsic dynamical systems*, stable-Lipschitz SCMs are closed under a number of structural transformations (see Appendix 9.3 for more details). In particular, they are closed under the twin operation:

Theorem 7 (Closure under Twin Operation)**.** *Let M be stable-Lipschitz. Then M^{twin} is also stable-Lipschitz.*

Conjecture 1 (Counterfactual dGMP)**.** *Let M be stable-Lipschtiz. Then the counterfactual distributions of M satisfy the directed global Markov property relative to the corresponding twin network.*

I believe Conjecture 1 holds if the condition of independent noise in Theorem 5 can weakened, as an immediate consequence of Theorems 6 and 7. However, I would place 2:1 odds that I'm missing some additional aspect of the proof.

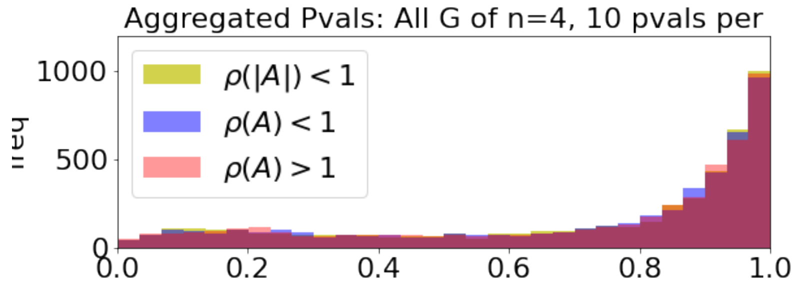

Figure 10: Aggregated pvals of conditional independence tests, across multiple SCMs. While each SCM was assessed individually, this plot clearly highlights the nigh-identical behavior of the three subclasses of Lipschitz SCMs that were tested. (I suspect the particular shape of the distribution may highlight something about fcit producing more confident independece predictions on certain graph structures (which were sampled equally among all three strata).)

5. Numerics

When I discovered a mistake in my attempt at proving Theorem 5, I sought to verify my conjecture numerically first, to guide my next research steps. Specifically, I wanted to test my belief that stable-Lipschitz SCMs satisfy d-separation, and that non-stable Lipschitz SCMs do not. Formally,

Conjecture 2 (stable-Lipschitz)**.** *Every stable-Lipschitz SCM satisfies the dGMP; that is, the observational distribution respects every conditional independence in the causal graph.*

Conjecture 3 (Lipschitz)**.** *Lipschitz SCMs don't generally satisfy the dGMP.*

The methodology I followed for my numerics is as follows:

- Generate all cyclic graphs of 4 nodes (totalling 107 non-isomorphic graphs)
- Given a graph, sample a neural network with Relu activations which respects that graph (and which has a spectral radius appropriate to the class of SCMs being tested)
- Enumerate all possible d-separations for this graph
- For each independence claimed by d-separation, sample a dataset from the observational distribution 10 times (by means of sampling from $P(U)$, then applying root-finding methods to avoid numerical instability)
- Test whether the claimed independence from d-separation holds in the observation distribution using `fcit`

The surprising results of the numerics are in Figure 10. Conjecture 2 held (and indeed I was able to prove Theorem 5 this week), but Conjecture 3 appears to be refuted! That is, it would appear that not only does d-separation hold for stable-Lipschitz SCMs, it holds for general Lipschitz SCMs as well! It will be interesting to see if the proof can be extended to Lipschitz SCMs.

6. Future Work

This research only scratches the surface of an area previously seen as intractable. Consequently, I think there may be a lot of low-hanging fruit.

6.1. Pearl Causal Hierarchy

[1] sets up a nice framework for the PCH, and suggest that the authors believe it holds, but stops short of actually making any claims about 'collapse on a set of measure zero'.

To highlight why proving the PCH seems nontrivial to me, I invite you to imagine what a negative result might look like:

- Even though the PCH holds for acyclic SCMs,
- since acyclic SCMs are themselves a set of measure 0 among the space of cyclic SCMs,
- it turns out that non-collapse is in fact the exception rather than the norm.

Nevertheless, I suspect the PCH will indeed hold for cyclic SCMs. But this should be demonstrated before spending lots of time on do-calculus or cyclic counterfactuals, for example, since if the PCH collapses then identification becomes much easier.

6.2. *Extending Observational dGMP*

As discussed following Theorem 5, I believe the constraints of independent and additive noise can be weakened. Furthermore, the numerics suggest that even the assumption of simple-Lipschitz may be stronger than is necessary.

6.3. *Do-calculus*

The do-calculus would generalize the backdoor result of this report.

I suspect that the do-calculus should already follow from the current results of this report; however, I would like to take a fine-tooth comb through the validity proof of the acyclic do-calculus first, to ensure that everything generalizes properly to cyclic SCMs.

6.4. *Counterfactual dGMP*

Since the dGMP has been proven to hold for both the observational and interventional distributions of simple-Lipschitz SCMs, it is natural to ask if it also holds for counterfactual distributions:

Problem 3 (stable-Lipschitz SCMs and d-separation: Counterfactual)**.** *Prove whether the counterfactual distribution of stable-Lipschitz SCMs satisfy the directed global Markov property: that is, whether every conditional independence read-off by d-separation in the twin network holds in the corresponding counterfactual distribution of the twin SCM.*

6.5. *Multiple Equilibria*

I conjecture that the results of this report may generalize nicely to non-uniquely-solvable SCMs, *providing a general form of causal identification which does not care which equilibria is generating the potential response!*

This is the area of research I am most excited by, because currently the multiplicity of equlibria is a signicant difficulty for analysis of game theory and economics. Indeed, this is the fundamental premise of the settable systems framework [14].

Settable systems takes an explicit approach toward modeling multiple equilibira. I conjecture that an implicit approach may be sufficient. In this way, a significant amount of additional modeling machinery of settable systems may be set aside.

My overall strategy for this direction basically consists of weakening the "uniquely solvable" condition to "solvable", by showing that d-separation holds in the neighborhood of equilibria of "locally stable-Lipschitz" SCMs. Intuitively I think this should work because in some sense, knowing which equilibria we're at is SCM-level knowledge – more than we should need for identification from the causal graph G and $P(V)$.

7. Conclusion

Despite being the natural modeling choice for phenomena involving feedback, cyclic causality has seen limited use because many of the convienent guarentees of acylic SCMs fail to hold for general cyclic SCMs. In particular, d-separation on the causal graph loses validity in the general cyclic setting!

In this report, we have presented research about the validity of d-separation (aka. the dGMP: directed global Markov property) for cyclic SCMs. Specifically, we proposed the class of stable-Lipschitz SCMs, which generalize acyclic SCMs and are contained by simple SCMs, thus inheriting all of the nice properties of simple SCMs. Stable-Lipschitz SCMs behave like linear SCMs asyptotically, and indeed, are proven to satsify the dGMP (currently with some additional assumptions). These results are verified numerically, and the validity of the backdoor criterion is proven as well. We also show that stable-Lipschitz SCMs are closed under interventions, and that the interventional distributions of simple SCMs are similarly shown to satisfy the dGMP.

Lastly, we discussed further directions for research, including the Pearl Causal Hierarchy, the do-calculus, and most exciting, the possibility of causal identification with multiple equilibria.

8. Acknowledgements

This project is being submitted jointly for Elias Bareinboim's Causal Inference II course, and David Blei's Applied Causality course. I am indebted to both of their feedback and advice, as well as Alexis Bellot, my assigned mentor for CI2.

References

1. Stephan Bongers, Patrick Forré, Jonas Peters, and Joris M. Mooij. Foundations of structural causal models with cycles and latent variables. *The Annals of Statistics*, 49(5):2885 – 2915, 2021.
2. Leonid Bunimovich, D. J. Passey, Dallas Smith, and Benjamin Webb. Spectral and dynamic consequences of network specialization. *International Journal of Bifurcation and Chaos*, 30(06):2050091, may 2020.

3. Leonid Bunimovich, Dallas Smith, and Benjamin Webb. Finding hidden structures, hierarchies, and cores in networks via isospectral reduction. *Applied Mathematics and Nonlinear Sciences*, 4(1):231–254, 2019.
4. Leonid Bunimovich and Benjamin Webb. *Isospectral Transformations: A New Approach to Analyzing Multidimensional Systems and Networks*. Springer, 2014.
5. Camille Carter, Jacob Murri, David Reber, and Benjamin Webb. Stability of stochastically switched and stochastically time-delayed systems, 2020.
6. Krzysztof Chalupka, Pietro Perona, and Frederick Eberhardt. Fast conditional independence test for vector variables with large sample sizes, 2018.
7. Patrick Forr'e and Joris M. Mooij. Markov properties for graphical models with cycles and latent variables. *arXiv: Statistics Theory*, 2017.
8. Lewis Hammond, James Fox, Tom Everitt, Ryan Carey, Alessandro Abate, and Michael Wooldridge. Reasoning about causality in games. *Artificial Intelligence (submitted)*, 2022.
9. R.A. Horn and C.R. Johnson. *Matrix Analysis*. Cambridge University Press, 2012.
10. Antti Hyttinen, Frederick Eberhardt, and Patrik O. Hoyer. Learning linear cyclic causal models with latent variables. *J. Mach. Learn. Res.*, 13:3387–3439, 2012.
11. Judea Pearl. *Causality*. Cambridge University Press, 2 edition, 2009.
12. David Reber and Benjamin Z. Webb. Intrinsic stability: Global stability of dynamical networks and switched systems resilient to any type of time-delays. *ArXiv*, abs/1906.04000, 2019.
13. Peter Spirtes. Directed cyclic graphical representations of feedback models. *CoRR*, abs/1302.4982, 2013.
14. Halbert White and Karim Chalak. Settable systems: An extension of pearl's causal model with optimization, equilibrium, and learning. *Journal of Machine Learning Research*, 10(61):1759–1799, 2009.
15. Halbert White, Haiqing Xu, and Karim Chalak. Causal discourse in a game of incomplete information. *Journal of Econometrics*, 182(1):45–58, 2014. Causality, Prediction, and Specification Analysis: Recent Advances and Future Directions.

9. Appendix

9.1. Notation

- Lowercase letters (e.g. z) for particular values of random variables (Z); bold for sets of variables ($\mathbf{z}$ and $\mathbf{Z}$).
- $\text{dom}(X)$ denotes the domain over which the random variable X is defined.
- $\sigma(A)$: The eigenvalues of a square matrix A.
- $\rho(A)$: the spectral radius of A.
- $An(\mathbf{Z})$: the ancestors of $\mathbf{Z}$ in a causal diagram, including $\mathbf{Z}$ itself.
- For matrices A and B, $A \preceq B$ means that $a_{ij} \leq b_{ij}$ for all i, j.

9.2. Proofs

Proof of Theorem 1: acylic $\subset$ stable-Lipschitz. Let G be the causal diagram of M. G has no strongly connected components, so the conditions for stable-Lipschitz are vacuously satisfied. □

Proof of Theorem 4: Closure under Interventions. Let G be the causal diagram of M.

Case 1: G is acyclic. Then $G_{\overline{\mathbf{X}}}$ is also acyclic, so vacuously $M_{\text{do}(\mathbf{X}=\mathbf{x})}$ is stable-Lipschitz.

Case 2: G is strongly connected. Let A be the Lipschitz matrix of F. Since M is stable-Lipschitz, $\rho(A) < 1$. Let B be defined element-wise as

$$b_{ij} = \begin{cases} 0 & i, j \in \mathbf{X} \\ a_{ij} & \text{otherwise} \end{cases}$$

B is the Lipschitz matrix for $M_{\text{do}(\mathbf{X}=\mathbf{x})}$, because $\text{do}(\mathbf{X} = \mathbf{x})$ sets each component function in $\mathbf{X}$ to a constant, inducing everywhere-zero partial derivatives. Furthermore, $0 \preceq B \preceq A$, so $\rho(B) \leq \rho(A) < 1$. Thus, $M_{\text{do}(\mathbf{X}=\mathbf{x})}$ is stable-Lipschitz.

Case 3: G contains multiple strongly connected components. Let $\{\mathbf{Z}_s\}$ be the set of the strongly connected components of G, and $\{\mathbf{W}_t\}$ the set of the set of the strongly connected components of $G_{\overline{\mathbf{X}}}$. Since $G_{\overline{\mathbf{X}}}$ is a subgraph of G, we have that for each $\mathbf{W}_t$ there exists a $\mathbf{Z}_s$ such that $\mathbf{W}_t \subseteq \mathbf{Z}_s$. Since $F_{\mathbf{Z}_s}$ is differentiable and Lipschitz by hypothesis, so is $F_{\mathbf{W}_t}$. As in Case 2, let $A_{\mathbf{Z}_s}$ be the Lipschitz matrix of $F_{\mathbf{Z}_s}$, which by hypothesis satisfies $\rho(A_{\mathbf{Z}_s}) < 1$. Let B be defined over $\mathbf{Z}_s$ as

$$b_{ij} = \begin{cases} a_{ij} & i, j \in \mathbf{W}_t \\ 0 & \text{otherwise} \end{cases}$$

Clearly $0 \preceq B \preceq A_{\mathbf{Z}_s}$, so $\rho(B) \leq \rho(A_{\mathbf{Z}_s}) < 1$. Since the principle submatrix $B_{\mathbf{W}_t}$ is precisely the Lipschitz matrix of $F_{\mathbf{W}_t}$, we have $\sigma(B_{\mathbf{W}_t}) \subseteq \sigma(B)$. Thus $\rho(B_{\mathbf{W}_t}) < 1$. Since $\mathbf{W}_t$ was an arbitrary strongly connected component, $M_{\mathrm{do}(\mathbf{X}=\mathbf{x})}$ is stable-Lipschitz. □

Proof of Lemma 2: Unique Solvability. Let G be the causal diagram of M.

Case 1: G is acyclic. It was proven in [1] that M is uniquely solvable if it is acyclic.

Case 2: G is strongly connected. Let $\mathbf{u} \in \mathrm{dom}(\mathbf{U})$. By [12], there exists a unique, globally attracting fixed point $\mathbf{v}^* \in \mathrm{dom}(\mathbf{V})$ of the dynamical system $\left(F_{\mathbf{U}=\mathbf{u}}(\mathbf{V}), \mathbf{V}\right)$; that is, for all $\mathbf{v}_0 \in \mathrm{dom}(\mathbf{V})$, $\lim_{k\to\infty} F^k_{\mathbf{U}=\mathbf{u}}(\mathbf{v}_0) = \mathbf{v}^*$. Thus the structural equations $\mathbf{V} = F(\mathbf{V}, \mathbf{u})$ have the unque solution $\mathbf{v}^*$, so M is uniquely solvable.

Case 3: G is not strongly connected. Let $\{\mathbf{Z}_i\}_{i=1}^s$ be the set of strongly connected components of G in topological order. Let us partition the vertices not contained in strongly connected components, based on the order in which they impact a strongly connected component. Specifically, let $\{\mathbf{W}_i\}_{i=1}^{s+1}$ be defined as

$$\begin{aligned}
\mathbf{W}_1 &= An(\mathbf{Z}_1) \setminus \mathbf{Z}_1 \\
\mathbf{W}_2 &= An(\mathbf{Z}_2) \setminus (\mathbf{Z}_1 \cup \mathbf{Z}_2 \cup \mathbf{W}_1) \\
&\vdots \\
\mathbf{W}_s &= An(\mathbf{Z}_s) \setminus \left((\cup_{i=1}^s \mathbf{Z}_i) \cup (\cup_{i=1}^{s-1} \mathbf{W}_i)\right) \\
\mathbf{W}_{s+1} &= \mathbf{V} \setminus \left((\cup_{i=1}^s \mathbf{Z}_i) \cup (\cup_{i=1}^{s} \mathbf{W}_i)\right)
\end{aligned}$$

where we may have a $\mathbf{W}_i = \emptyset$. By construction, $\{\mathbf{Z}_i\}_{i=1}^s \cup \{\mathbf{W}_i\}_{i=1}^{s+1}$ forms a partition of $\mathbf{V}$, and $(\mathbf{W}_1, \mathbf{Z}_1, \ldots, \mathbf{W}_s, \mathbf{Z}_s, \mathbf{W}_{s+1})$ is a valid topological ordering of G. Consequently, each $\{\mathbf{W}_1\}$, $\{\mathbf{W}_1, \mathbf{Z}_1\}$, $\{\mathbf{W}_1, \mathbf{Z}_1, \mathbf{W}_2\}$, etc. is ancestral in G.

We evaluate the solution of $\mathbf{V} = F(\mathbf{V}, \mathbf{u})$ iteratively through $(\mathbf{W}_1, \mathbf{Z}_1, \ldots, \mathbf{W}_s, \mathbf{Z}_s, \mathbf{W}_{s+1})$ by pushing fixed values through each $\mathbf{W}_i$ to obtain a unique fixed output; evaluating each $\mathbf{Z}_i$ with these fixed inputs to obtain the unique fixed point of the dynamical system; feeding these newly fixed values through $\mathbf{W}_{i+1}$, and so on. Specifically, because $\{\mathbf{W}_1\}$ is ancestral, we have $\mathbf{W}_1 = F_{\mathbf{W}_1}(\mathbf{V}, \mathbf{u}) = F_{\mathbf{W}_1}(\mathbf{W}_1, \mathbf{u})$ which is acyclic, and so has a unique fixed point $\mathbf{w}_1^*$. Next, $\mathbf{Z}_1 = F_{\mathbf{Z}_1}(\mathbf{V}, \mathbf{u}) = F_{\mathbf{Z}_1}(\mathbf{w}_1^*, \mathbf{Z}_1, \mathbf{u})$ has $\rho(A_{\mathbf{Z}_1}) < 1$, so by Case 2 produces a unique fixed point $\mathbf{z}_1^*$. Next, $\mathbf{W}_2 = F_{\mathbf{W}_2}(\mathbf{V}, \mathbf{u}) = F_{\mathbf{W}_2}(\mathbf{w}_1^*, \mathbf{z}_1^*, \mathbf{W}_2, \mathbf{u})$ is again acyclic, so produces a unique fixed point $\mathbf{w}_2^*$, and so on.

In this way we obtain a unique fixed point $\mathbf{v}^*$ for $\mathbf{V} = F(\mathbf{V}, \mathbf{u})$ equal by stacking $\mathbf{v}^* = [\mathbf{w}_1^*, \mathbf{z}_1^*, \ldots, \mathbf{w}_s^*, \mathbf{z}_s^*, \mathbf{w}_{s+1}^*]$. Thus M is uniquely solvable. □

Proof of Theorem 3: stable-Lipschitz ⊂ simple. Let $\mathbf{u} \in \mathrm{dom}(\mathbf{U})$, $\mathbf{Z} \subseteq \mathbf{V}$, and $\mathbf{w} \in \mathrm{dom}(\mathbf{V} \setminus \mathbf{Z})$. By Theorem 4, $M_{\mathrm{do}(\mathbf{w})}$ is stable-Lipschitz, and hence uniquely solvable by Lemma 2. This implies that there is a unique solution $\mathbf{v}^*$ to the equations $\mathbf{V} = F_{\mathrm{do}(\mathbf{w})}(\mathbf{V}, \mathbf{u})$. By definition of $\mathrm{do}(\mathbf{w})$ this is

$$\begin{aligned}
\mathbf{Z} &= F_{\mathbf{Z}}(\mathbf{V}, \mathbf{u}) = F_{\mathbf{Z}}(\mathbf{Z}, \mathbf{W}, \mathbf{u}) \\
\mathbf{W} &= F_{\mathbf{W}}(\mathbf{V}, \mathbf{u}) = \mathbf{w}
\end{aligned}$$

with the latter equation already solved. Plugging this into the first equation, we obtain $\mathbf{Z} = F_{\mathbf{Z}}(\mathbf{Z}, \mathbf{w}, \mathbf{u})$, which must be uniquely solvable since we know $\mathbf{V} = F_{\mathrm{do}(\mathbf{w})}(\mathbf{V}, \mathbf{u})$ is uniquely solvable.

Since $\mathbf{u} \in \mathrm{dom}(\mathbf{U})$ and $\mathbf{w} \in \mathrm{dom}(\mathbf{V} \setminus \mathbf{Z})$ were arbitrary, we have that M is uniquely solvable w.r.t. $\mathbf{Z}$. Since $\mathbf{Z} \subseteq \mathbf{V}$ was arbitrary, M is simple. □

Lemma 8 (Injectivity)**.** *Let M have structural functions of the form $F(\mathbf{V}, \mathbf{U}) = H(\mathbf{V}) + \mathbf{U}$ (additive noise). Furthermore, assume that for every ancestral $\mathbf{W} \subseteq \mathbf{V}$,*

- *$F_{\mathbf{W}}$ is Lipschitz*
- *$det(I_{|\mathbf{W}|} - A_{\mathbf{W}}) \neq 0$ (uniquely solvable)*

Then $h_{\mathbf{V}} := \mathbf{V} - H(Pa(\mathbf{V}))$ is injective for ancestral $\mathbf{W} \subseteq \mathbf{V}$.

Proof. Assume $h_{\mathbf{W}}(\mathbf{w}^{(1)}) = h_{\mathbf{W}}(\mathbf{w}^{(2)})$. This implies $\mathbf{w}^{(1)} - \mathbf{w}^{(2)} = K(\mathbf{w}^{(1)}) - K(\mathbf{w}^{(1)})$, so if we take the element-wise absolute value we have the element-wise inequality $|\mathbf{w}^{(1)} - \mathbf{w}^{(2)}| = |K(\mathbf{w}^{(1)}) - K(\mathbf{w}^{(1)})| \preceq A_{\mathbf{W}}|\mathbf{w}^{(1)} - \mathbf{w}^{(2)}|$. Thus $(I_{|\mathbf{W}|} - A_{\mathbf{W}})|\mathbf{w}^{(1)} - \mathbf{w}^{(2)}| \preceq 0$. By hypothesis $\det(I_{|\mathbf{W}|} - A_{\mathbf{W}}) \neq 0$ so $(I_{|\mathbf{W}|} - A_{\mathbf{W}})^{-1}$ exists, so in fact $|\mathbf{w}^{(1)} - \mathbf{w}^{(2)}| \preceq 0$. Thus $\mathbf{w}^{(1)} = \mathbf{w}^{(2)}$. □

Proof of Theorem 5: Observational dGMP. We show that M satisfies the SEPwared property of [7] (Definition 3.8.14). Let $h_{\mathbf{V}} := \mathbf{V} - H(\text{Pa}(\mathbf{V}))$ leading to $h_{\mathbf{W}}(\mathbf{w}) = \mathbf{U}_{\mathbf{W}}$ for ancestral $\mathbf{W} \subseteq \mathbf{V}$. By Lemma 8, $h_{\mathbf{W}}$ is bijective. Thus we can evaluate

$$|h'_{\mathbf{W}}(\mathbf{w})| = |\frac{dh_{\mathbf{W}}}{d\mathbf{w}}(\mathbf{w})| = |\det(I_{|\mathbf{W}|} - \frac{dH_{\mathbf{W}}}{d\mathbf{w}}(\mathbf{w}))|$$

By similar reasoning to the proof of Theorem 4, $\rho(A_{\mathbf{W}}) < 1$, and by definition $|\frac{dH_{\mathbf{W}}}{d\mathbf{w}}(\mathbf{w})| \preceq A_{\mathbf{W}}$, so by Theorem 8.1.18 of [9] we have that $\rho(\frac{dH_{\mathbf{W}}}{d\mathbf{w}}(\mathbf{w})) < 1$ for all $\mathbf{v} \in \mathbf{V}$. This constitutes a sufficient condition for

$$|h'_{\mathbf{W}}(\mathbf{w})| = |\det(I_{|\mathbf{W}|} - \frac{dH_{\mathbf{W}}}{d\mathbf{w}}(\mathbf{w}))| \neq 0$$

for all $\mathbf{v} \in \mathbf{V}$. Thus, SEPwared holds, so by [7] M satisfies the dGMP. $\square$

Proof of Corollary 1: Adjustment Formula. The proof of the BDC relies on 1. a d-separation condition, 2. the directed global Markov property, and 3. axioms of probability. So long as we use cyclic d-separation, 1 can be checked. Theorem 5 provides 2, and 3 is unaffected. $\square$

Proof of Theorem 6: Interventional dGMP. By Theorem 3, we have that M is simple. By [1], this implies that $G_{\overline{\mathbf{X}}} = G_{M_{\text{do}(\mathbf{X}=\mathbf{x})}}$. By Theorem 4, we have that $M_{\text{do}(\mathbf{X}=\mathbf{x})}$ is also stable-Lipschitz.

All that remains is to show that the hypothesis of Theorem 5 still holds. Interventions do not change the structural form of F, so the additive noise condition still holds; similarly, independant noise is unaffected. Lastly, if $P_M(\mathbf{V})$ has density, so does $P_{M_{\text{do}(\mathbf{X}=\mathbf{x})}}(\mathbf{V})$. Thus M satisfies the dGMP. $\square$

Proof of Theorem 7: Closure under Twin Operation. By hypothesis, for each strongly connected component the corresponding $A_{\mathbf{Z}}$ satisfies $\rho(A_{\mathbf{Z}}) < 1$. Note that the strongly connected components of M^{twin} are as follows:

$$\text{strong}(F^{\text{twin}}) = \text{strong}(F) \cup \text{strong}(F')$$

that is, simply duplicated. Hence $A_{\mathbf{Z}}$ is the Lipschtiz matrix for both the original component and its duplicate, so indeed M^{twin} is also stable-Lipschitz. $\square$

9.3. *Intrinsic Dynamical Networks*

Considered from a dynamical-systems perspective, intrinsically stable dynamical networks (the theoretical foundation for stable-Lipschitz SCMs) have a number of promising properties relevant to causality. They have a unique equilibrium, so the potential response function of the SCM will be well-defined. They 'behave' like linear systems: they are asymptotically bounded by the dynamics of a linear system, and share the same equilibria if subject to the same forcing factor (exogenous distribution). Intuitively speaking, this means that we can go to linear-world, prove things about the system there, and the results will still hold in nonlinear-world. Furthermore, intrinsically-stable systems are quite general: they only require Lipschitz-continuity, and the domain to be a product of metric spaces (which could be something nice like $\mathbb{R}^4$, or something abstract like language and shapes). If a general domain like metric spaces is used, the linearization uses the metrics to map to the real numbers: in this way the domain is simplified as well.

Of particular interest for causality, intrinsically-stable systems derive their name because they are closed under a surprising number of structural transformations; that is, the resulting system will still be intrinsically-stable, and often (if applicable) the equilibria will be preserved. Some especially relevant transformations are: lengthening of paths (i.e. through time-delays [12]), collapsing portions of the graph [3], duplicating portions of the graph (specialization [2]), time-varying structural switching [5], and any and all isospectral transformations (transformations which preserve the eigenvalues of the system [4]).